

From Search to Generative AI: Investor Information Acquisition Around Earnings *

Shuqing Huang[†] Han Yan[‡]

May 20, 2026

Abstract

Generative AI (GenAI) platforms are increasingly used to interpret public financial disclosures, potentially reshaping investor information acquisition. Using daily web-traffic data for leading GenAI platforms and major search engines, we document a pronounced reallocation of investor attention toward GenAI on earnings-announcement days, while traditional search traffic shows no comparable response. This shift is stronger when announcing firms are more complex, earnings surprises are larger, and retail participation is higher. GenAI platforms also generate disproportionately more referral traffic to free financial websites during earnings periods, consistent with investors using GenAI as a gateway for financial information. At the firm level, greater GenAI activity during earnings periods is associated with lower illiquidity and narrower bid-ask spreads, concentrated among high-retail-attention firms. These effects are specific to earnings announcements and attenuate after Google integrates synthesis into its search interface. Overall, the evidence suggests GenAI is broadening access to interpreted financial information around earnings disclosures.

*We are grateful for thoughtful comments and suggestions from Mary Billing, Elizabeth Blankespoor, Russell Lundholm and Dawn Matsumoto. We also thank conference participants at the 2026 Southern California Accounting Research Forum (SCARF) for their valuable feedback. We gratefully acknowledge financial support from California State University, Fullerton and University of British Columbia.

[†]California State University Fullerton, College of Business and Economics

[‡]University of British Columbia, Sauder School of Business

1 Introduction

A central premise of securities regulation is that public disclosure reduces information asymmetry between firms and investors. Yet public availability does not imply equal usability. Investors must first become aware that information exists, acquire the relevant data, and then integrate disparate pieces of information into valuation-relevant beliefs ([Blankespoor et al., 2020](#)). This final integration step is especially demanding because it requires users to combine quantitative results, management narratives, peer information, macroeconomic context, and market reactions into a coherent interpretation. These costs are unlikely to be borne equally across investors. Sophisticated institutions often have analysts, data infrastructure, proprietary models, and direct access to information intermediaries. Retail investors, by contrast, generally face greater attention, acquisition, and processing constraints, even when the underlying disclosures are publicly available. Prior research shows that individual investors often disregard accounting information because of awareness and acquisition costs ([Blankespoor et al., 2019](#)), and the SEC has long emphasized investor protection, fair markets, and the need for disclosure that ordinary investors can understand ([SEC, 1998, 2021, 2023](#)).

This availability-usability distinction is particularly important around earnings announcements. Earnings releases and conference calls concentrate dense numerical and narrative information into narrow windows, and investors must rapidly determine whether the news changes firm value. Prior work documents that investors respond to these constraints by reallocating attention across firms ([Hirshleifer et al., 2009](#); [Ferracuti and Lind, 2025](#)) and by using digital search tools to acquire information around earnings announcements ([Da et al., 2011](#); [Drake et al., 2012](#)). These studies implicitly assume that search engines are the dominant online gateway through which investors seek financial information. That assumption has become less certain. Investors now have access to generative artificial intelligence (GenAI) platforms that do more than retrieve documents. They can summarize earnings releases, explain conference call discussions, compare

a firm's performance with peers, and translate complex disclosure into conversational responses. This technological shift raises a question that is both empirically and policy relevant: does GenAI reduce the processing frictions that prevent broader investor participation in the use of public disclosures, or does it create a new technological divide between investors who can use AI effectively and those who cannot?

The answer is *ex ante* ambiguous. On one hand, GenAI may help level the playing field. By converting complex disclosure into synthesized, plain-language interpretation, GenAI can reduce the integration costs that are especially burdensome for retail investors. This possibility is consistent with a long-standing regulatory objective: disclosure is most valuable when investors can actually use it. The SEC's plain-English initiative recognized that disclosure documents must be understandable, not merely available ([SEC, 1998](#)). The XBRL mandate similarly reflected the idea that technology can reduce information frictions and make financial data more accessible and comparable, with prior research showing that machine-readable reporting can reduce disadvantages faced by smaller market participants ([Hodge et al., 2004](#); [Bhattacharya et al., 2018](#)). GenAI may represent the next step in this evolution.

On the other hand, GenAI may not democratize information processing. Its benefits may depend on investors' awareness of the tools, access to them, and ability to prompt and verify AI-generated output. Some investors may not know which AI tools are available or may use them only as a more convenient search box. Others may distrust AI-generated answers, prefer direct links to primary sources, or seek only a simple factual lookup for which traditional search is sufficient. Even among users who adopt GenAI, the value of the response depends on the user's ability to ask the right question and evaluate the answer. Sophisticated investors may therefore benefit more because they have better complementary information, stronger incentives to experiment, greater expertise to detect hallucinations, and, in some cases, in-house AI infrastructure. These considerations echo broader regulatory concerns about technology-mediated investing.

Digital platforms can reduce participation costs, but they can also shape investor attention, increase overconfidence, or create conflicts through choice architecture and predictive analytics ([Barber and Odean, 2000](#); [Barber et al., 2022](#); [SEC, 2021, 2023](#)). Thus, whether GenAI narrows or widens informational advantages remains unclear.

To examine this question, we study how GenAI affects investor information acquisition around earnings announcements. The economic premise is that traditional search engines and GenAI platforms serve qualitatively different functions in the disclosure-processing chain. Search engines retrieve documents that match a query; they reduce awareness and acquisition costs but largely leave the burden of synthesis to the user. GenAI platforms, by contrast, generate contextualized responses that combine information across sources; they operate more directly on the integration margin. Importantly, the two technologies are not mutually exclusive. Investors may use GenAI to obtain an initial synthesis and then use search to verify details, inspect source documents, or obtain raw data. This distinction motivates our analysis of how investors allocate attention across retrieval-oriented and synthesis-oriented tools when earnings-information-processing demands rise.

We test these ideas using daily web-traffic data from SimilarWeb for ten leading GenAI platforms and five traditional search engines over the period August 2024 through October 2025. We measure earnings-announcement intensity using value-weighted counts of firms announcing earnings, constructed from I/B/E/S and Compustat. Our baseline research design compares the traffic response of GenAI platforms to that of traditional search engines on days with different levels of earnings-announcement intensity. The design includes website fixed effects, time fixed effects, and controls for platform user characteristics, thereby isolating whether GenAI platforms experience a differential increase precisely when the market faces a larger earnings-information-processing burden. We then examine when this response is strongest and whether it is associated with changes in the broader information environment. Specifically, we study whether the GenAI response is concentrated in high-interpretation-demand and retail-attention settings,

whether it is linked to downstream visits to financial information websites, and whether it is associated with market outcomes that reflect information asymmetry and price discovery.

We report four main findings. First, GenAI platforms experience a disproportionate increase in web traffic on days with greater earnings-announcement intensity. In our tightest specification, which absorbs common daily variation through date fixed effects, moving from the lowest to the highest decile of earnings-announcement intensity is associated with approximately a 14 percent larger increase in traffic to GenAI platforms relative to traditional search engines. At the same time, traditional search traffic does not increase on high-announcement days. This null response contrasts with evidence from prior work showing that Google search activity rises around earnings announcements and serves as a revealed measure of investor attention ([Da et al., 2011](#); [Drake et al., 2012](#)). The evidence suggests that earnings-related information acquisition has partly migrated from retrieval-oriented search toward synthesis-oriented GenAI.

Second, the reallocation toward GenAI is strongest in settings where interpretation demands are high and where retail investors are most active. The differential GenAI response is concentrated on days when announcing firms are more complex and earnings surprises are greater, consistent with GenAI's synthesis function being most valuable when earnings information is harder to interpret. The response is also significantly stronger on days when announcing firms have high retail investor participation, indicating that GenAI adoption around earnings is especially pronounced in the market segment where broad access to interpreted information matters most.

Third, GenAI platforms generate disproportionately more referral traffic to free financial websites during high-earnings-announcement periods. This result provides a mechanism linking the platform-level traffic response to financially relevant information acquisition. It is also consistent with a complementary information process: users may first obtain an AI-generated synthesis and then visit accessible financial portals, such as

Yahoo Finance or Seeking Alpha, to verify facts, examine raw data, or obtain additional context. The concentration of referrals to free financial destinations is also consistent with broad investor use because these sources are more accessible to nonprofessional investors than paywalled institutional platforms.

Fourth, GenAI activity during high-announcement periods is associated with improved market outcomes in the firms for which earnings information is most likely to require interpretation. Specifically, GenAI traffic is associated with lower illiquidity and narrower bid-ask spreads for firms with high retail attention, high social media attention, and longer conference call transcripts. Traditional search traffic does not exhibit the same pattern. These market-outcome results speak directly to the *ex ante* ambiguity in the paper. If GenAI primarily strengthens sophisticated investors, one might expect greater adverse selection and wider spreads. If GenAI helps a broader investor base process public disclosure more quickly, information asymmetry should fall and liquidity should improve. The evidence is more consistent with the latter interpretation, while we remain cautious in drawing welfare conclusions because we do not observe the identity of individual GenAI users or the quality of their decisions.

We conduct several additional tests to address alternative explanations. Placebo tests using non-earnings corporate events such as M&A and legal actions show that the traffic reallocation is specific to earnings announcements rather than a general response to corporate news. The baseline findings are also robust to alternative measures of web traffic, alternative measures of earnings-announcement intensity, and timing adjustments for after-market announcements. Finally, we exploit Google's introduction of AI Mode, which embedded generative synthesis capabilities into Google Search. After this convergence shock, the differential traffic response of GenAI platforms attenuates, consistent with the interpretation that the baseline effect is tied to synthesis functionality rather than unrelated platform characteristics.

This paper contributes to three streams of research. First, it extends the literature

on investor attention and information acquisition. Search activity has long served as a primary digital trace of investor information demand, especially for retail investors (Da et al., 2011; Drake et al., 2012; Barber and Odean, 2008). We show that in the GenAI era, search activity may understate true earnings-related information acquisition because part of the information journey has shifted toward synthesis-oriented platforms. Second, the paper contributes to research on disclosure processing costs and retail investor frictions. Prior work shows that technologies such as XBRL and search-facilitating tools can reduce some information frictions, but retail investors may still fail to use public accounting information when awareness, acquisition, or processing costs remain high (Hodge et al., 2004; Blankespoor et al., 2019; Bhattacharya et al., 2018; Blankespoor et al., 2020). We provide field evidence that a new technology operates more directly on the integration margin and is associated with improvements in the information environment in retail-relevant settings. Third, the paper contributes to the emerging literature on GenAI and capital markets. Recent studies show that GenAI affects trading, analyst forecasting, and the production of financial information (Cheng et al., 2025; Bertomeu et al., 2025; Blankespoor et al., 2026a,b). We add to this literature by documenting how investor information acquisition reallocates between existing search technologies and new synthesis technologies during a central disclosure event, and by examining whether the associated improvement in the information environment reflects broader investor participation or concentrated sophistication.

Our findings also speak to ongoing policy debates over technology, investor protection, and the democratization of financial markets. Regulators and policymakers have repeatedly faced the tension that greater access can benefit ordinary investors but can also expose them to new risks if information is complex, intermediaries shape behavior, or investors cannot evaluate the quality of what they receive. This tension appears in debates over plain-English disclosure, XBRL, digital engagement practices, predictive analytics, and broader access to private markets (SEC, 1998, 2021, 2023; Jackson, 2019). GenAI adds a

new dimension to this debate. It may function as a user-side interpretation technology that makes public disclosure more usable, but it may also introduce new verification burdens and uneven benefits across investor types. By showing when investors shift from search to GenAI and how this shift relates to liquidity and spreads, our paper provides early evidence on how AI-assisted interpretation is reshaping the public-disclosure information environment.

The remainder of the paper proceeds as follows. Section 2 describes the institutional background of GenAI platforms and their interaction with traditional search engines. Section 3 reviews the related literature. Section 4 describes the data and variable measurement. Section 5 presents the research design and empirical results. Section 6 reports robustness tests. Section 7 concludes.

2 Institutional Background

Generative artificial intelligence platforms emerged as widely accessible consumer tools beginning in late 2022, when OpenAI released ChatGPT. Within months, competing platforms—including Anthropic’s Claude, Google’s Gemini, and Perplexity AI—entered the market, each offering conversational interfaces capable of synthesizing information from heterogeneous sources in response to natural-language prompts. By mid-2025, ChatGPT alone attracted over 4.5 billion monthly visits, and the broader generative AI market was valued at approximately \$67 billion with annualized growth exceeding 24 percent (Liu et al., 2025). Unlike traditional keyword-based search engines, which index the web and return ranked lists of hyperlinks for the user to navigate, generative AI systems rely on large language models (LLMs) trained on extensive text corpora to produce direct, synthesized responses. The technological distinction is fundamental: search engines perform information retrieval, locating and ordering pre-existing documents that match user-supplied terms, whereas generative AI performs information synthesis,

constructing novel text that integrates content from multiple underlying sources into a coherent narrative ([Zubair et al., 2025](#)). This difference in architecture implies a corresponding difference in the cognitive demands placed on the user, because the labor of reading, filtering, and reconciling disparate documents—which traditional search delegates entirely to the human—is partially automated by the generative model.

The economic distinction between these two technologies maps onto the disclosure processing cost framework developed in [Blankespoor et al. \(2020\)](#). That framework decomposes the cost of using public information into three sequential components: awareness costs (learning that a disclosure exists), acquisition costs (obtaining the raw disclosure), and integration costs (converting acquired data into a valuation-relevant signal). Traditional search engines primarily reduce acquisition costs by lowering the effort required to locate relevant documents, but they leave integration costs largely unchanged; the user must still read, interpret, and reconcile the retrieved materials. Generative AI, by contrast, operates most directly on the integration margin. When an investor prompts a platform to explain an earnings surprise or summarize a conference call transcript, the model performs the synthesis that would otherwise require substantial cognitive effort. Consistent with this characterization, experimental evidence finds that users interacting with LLM-based tools complete information-intensive tasks faster than users relying on conventional web search, issuing fewer but more complex queries and producing decisions of comparable or higher accuracy ([Spatharioti et al., 2025](#)). At the same time, LLM-generated outputs introduce new frictions: users may face greater verification needs because responses can contain errors or hallucinated statements and because citations may be incomplete or difficult to evaluate ([Spatharioti et al., 2025](#); [Zubair et al., 2025](#)). In addition, [Melumad and Yun \(2025\)](#) show that LLM-assisted information acquisition leads to shallower learning structures and reduced recall relative to traditional web search, even while improving efficiency and user experience. These frictions carry distributional implications. The ability to craft effective prompts and

to critically evaluate AI-generated output depends on domain knowledge and financial sophistication, attributes that are unevenly distributed across investor types. Sophisticated investors, who possess deeper financial expertise and often have access to proprietary AI systems integrated into professional workflows, may extract substantially more value from GenAI than retail investors who use the technology as a simple question-answering tool without fully exploiting its synthesis capabilities.

The competitive interaction between generative AI and traditional search engines is better characterized as partial substitution than full displacement. Clickstream data from the second half of 2024 show that ChatGPT's referral traffic grew from fewer than 10,000 unique destination domains in July 2024 to over 30,000 by November 2024, indicating that generative AI increasingly serves as a starting point for information journeys that subsequently direct users to primary sources ([Semrush, 2025a](#)). At the same time, user intent differs markedly across the two technologies. Analysis of 80 million clickstream records reveals that only 30 percent of ChatGPT prompts align with the four canonical search-intent categories—navigational, informational, commercial, and transactional—that characterize traditional engine queries; the remaining 70 percent reflect novel interaction patterns oriented toward advice, interpretation, and content creation ([Semrush, 2025b](#)). These developments can also reshape the broader information ecosystem because they change users' incentives to search other websites. When an AI-generated summary adequately satisfies the user's informational need, users may be less likely to open external links, thereby concentrating information consumption within the AI interface ([Pew Research Center, 2025](#)). In sum, generative AI platforms and LLM-enhanced search engines are increasingly competing not only for user attention, but also for control over the initial gateway through which public information is accessed and interpreted.

3 Related Literature

This paper relates to three streams of literature. The first studies investor attention and information acquisition around corporate disclosures. A central insight from this literature is that investors have limited attention and therefore rely on information intermediaries and digital search tools to identify and process value-relevant news. [Hirshleifer et al. \(2009\)](#) show that investors underreact to earnings news when attention is distracted by competing announcements, and [Ferracuti and Lind \(2025\)](#) show that investors acquire more macroeconomic information around earnings clusters. [Da et al. \(2011\)](#) establish Google search activity as a revealed measure of investor attention, particularly retail attention, and [Drake et al. \(2012\)](#) show that search activity increases around earnings announcements and is associated with price discovery. Related work shows that digital platforms can shape which disclosures investors notice and how quickly they react to them (e.g., [Lawrence et al., 2018](#)). These studies largely treat search activity as the primary observable trace of online information demand. We extend this literature by studying a setting in which investors can choose between retrieval-oriented search engines and synthesis-oriented GenAI platforms. Our evidence suggests that the relevant margin is no longer only how much investors search, but also which technology they use to transform public disclosure into decision-relevant information.

The second stream examines disclosure processing costs, retail-investor frictions, and the role of technology in making public information usable. [Blankespoor et al. \(2020\)](#) organize the cost of using public disclosure into awareness, acquisition, and integration costs. Much prior work studies technologies that reduce the first two components. For example, [Hodge et al. \(2004\)](#) show that search-facilitating technology can improve nonprofessional users' ability to locate and integrate related financial statement information, and [Bhattacharya et al. \(2018\)](#) show that the SEC's XBRL mandate reduced informational disadvantages faced by smaller institutions. Yet public availability and technological access do not guarantee use. [Blankespoor et al. \(2019\)](#) show that individual

investors may disregard accounting information when awareness and acquisition costs remain high. Our paper builds on this work by studying a technology that potentially operates more directly on the integration margin. GenAI does not simply make a filing easier to find or a data item easier to extract; it can summarize, compare, and contextualize information. This distinction allows us to examine whether a new disclosure-processing technology is associated with broader information acquisition in settings where interpretation costs are high and retail attention is salient.

This retail-investor angle also connects our paper to research on financial intermediaries and the unequal distribution of information-processing capacity. Sell-side analysts, the media, and financial platforms help investors interpret disclosures, but their benefits may differ across investor groups. Prior work shows that individual investors respond differently from sophisticated investors to information signals and can be especially sensitive to attention, salience, and intermediary design ([Barber and Odean, 2008](#); [Barber et al., 2022](#)). Analysts and other intermediaries can reduce processing costs for less sophisticated investors, but sophisticated investors often have additional advantages, including proprietary data, internal research capacity, and direct access to management. GenAI may therefore have two competing effects. It may substitute for part of the interpretation function historically supplied by analysts and other intermediaries, thereby helping ordinary investors process complex public information more quickly. Alternatively, it may magnify the advantage of sophisticated investors if effective use requires domain knowledge, verification skill, prompt quality, or proprietary complementary information. We use this tension to interpret both our cross-sectional evidence in high-retail-attention settings and our market-outcome tests.

The third stream is the emerging literature on GenAI and capital markets. Recent studies show that GenAI affects investor behavior, information production, and market outcomes. [Cheng et al. \(2025\)](#) use ChatGPT outages to show that AI availability affects trading activity and price informativeness. [Bertomeu et al. \(2025\)](#) exploit Italy's temporary

ban of ChatGPT and find effects on analyst forecasts. [Blankespoor et al. \(2026a\)](#) study retail investors' use of GenAI for financial information processing, while [Croom \(2025\)](#) examines how interactive AI affects investor judgments. Other work studies how GenAI changes the production of financial information: [Bradshaw et al. \(2025\)](#) examine AI-authored articles on Seeking Alpha, and [Blankespoor et al. \(2026b\)](#) document firms' use of GenAI in financial reporting. [Ecker et al. \(2025\)](#) analyze stock-related queries on a Chinese GenAI platform and show that user-platform interaction can reveal how investors employ AI for market-related tasks. Collectively, this literature establishes that GenAI matters for capital markets, but most evidence studies AI adoption, AI availability, or AI-generated content in isolation. Less is known about whether GenAI changes the allocation of information acquisition across existing and new technologies during recurring public-disclosure events.

Our contribution is to connect these literatures by studying the migration of investor information acquisition from search to GenAI around earnings announcements. Earnings announcements provide a setting in which disclosure is public, time-sensitive, and interpretation-intensive. By comparing GenAI platforms with traditional search engines, we can distinguish a shift toward synthesis from a general increase in online attention. By examining retail-attention settings, referral traffic to free financial websites, and market liquidity outcomes, we also provide evidence relevant to whether the resulting information environment is more consistent with broader investor participation or concentrated sophistication. In doing so, the paper extends the investor-attention literature from a single-technology search paradigm to a multi-technology environment in which the choice of information-processing tool is itself economically meaningful.

4 Data and Measurement

4.1 Web Traffic and Platform Classification

We obtain digital traffic data from SimilarWeb, a digital intelligence platform that provides granular measures of website usage, visitor demographics, and referral pathways. SimilarWeb constructs its estimates by integrating four categories of signals: first-party analytics shared by participating websites, anonymized device-level behavioral data from a contributory network of consumer software products, partnership data from internet service providers and other data intermediaries, and automated public web data extraction. These heterogeneous inputs are cleaned, matched, and extrapolated using machine-learning-based procedures to produce statistically representative traffic estimates across countries, industries, and devices ([SimilarWeb, 2024c](#)). The platform reports coverage of over 100 million websites, with estimates derived from billions of digital signals processed daily ([SimilarWeb, 2024a,b](#)). SimilarWeb’s digital metrics are integrated into professional research environments—including the Bloomberg Terminal’s alternative data platform ([Bloomberg–SimilarWeb, 2026](#))—and have been used in prior academic research on digital information acquisition ([DeHaan et al., 2025](#)). These features make SimilarWeb well suited for measuring high-frequency shifts in user attention across information platforms.

Our primary dependent variable, *Web Traffic*, is the natural logarithm of total daily U.S. visits to each platform, constructed by aggregating desktop web and mobile application traffic. Combining both channels is important in our setting because user engagement with generative AI platforms and search engines increasingly spans devices; relying on either component alone would understate total observable interaction. We classify platforms into two groups based on the fundamental architectural distinction between information synthesis and information retrieval. The treatment group consists of ten generative AI platforms: ChatGPT, Perplexity, Claude, Gemini, Copilot, Grok, Deepseek, Meta AI, Poe,

and Qwen. These platforms accept natural-language prompts and return synthesized, conversational responses that integrate content from multiple sources. The control group consists of five traditional search engines: Google, Bing, Yahoo, DuckDuckGo, and Quora. These platforms index the web and return ranked lists of hyperlinks for the user to navigate. The indicator variable *GenAI* equals one for generative AI platforms and zero for traditional search engines. This classification is central to our identification strategy because the two platform types map onto different components of disclosure processing costs: search engines primarily reduce acquisition costs, whereas generative AI reduces integration costs (Blankespoor et al., 2020).

To account for compositional differences in user bases across platforms, we include six time-varying visitor-characteristic controls from SimilarWeb in all web-traffic regressions: the share of female visitors, the share of young visitors (aged 18–34), and four interest-based proxies—the fraction of visitors who also frequent investment-related websites (financial literacy), news and media websites (media attention), science and education websites (education), and social networking websites (social media engagement). These controls absorb variation in platform traffic that reflects demographic shifts or changes in the information diet of each platform’s user base rather than event-driven demand for financial disclosure processing.

SimilarWeb also provides monthly referral-traffic data that track visits originating from one website to another. We use these data to construct a measure of downstream financial information acquisition: *Referral Traffic*, defined as monthly referral visits from a source platform (either a generative AI platform or a traditional search engine) to financial destination websites (e.g., Yahoo Finance, Seeking Alpha, Bloomberg), scaled by the average daily web traffic of the source platform within the month and multiplied by 1,000. This scaling expresses referral intensity relative to the source platform’s overall traffic, allowing us to compare the extent to which different information platforms direct users to financial websites during earnings-announcement periods.

We note that SimilarWeb’s platform-level traffic data capture observable visits to websites and mobile applications and therefore do not include professional usage through APIs or indirect access via embedded integrations within third-party software. This limitation does not undermine our empirical design for two reasons. First, our focus is on general-public-accessible information acquisition around earnings announcements, where browser- and app-based interaction represents the primary channel for rapid, conversational interpretation of disclosures. Second, our identification strategy relies on relative reallocations of observable digital attention across platform types over time rather than on capturing the full universe of AI interactions, so unobserved API usage would bias our estimates only if it differentially shifts within the narrow earnings-announcement windows in a way that is systematically correlated with our treatment variable.

4.2 Earnings Announcement Intensity

We measure the intensity of earnings announcements on each calendar day using announcement dates from I/B/E/S. Following prior research that examines how the concentration of simultaneous earnings disclosures shapes investor attention and market outcomes ([Hirshleifer et al., 2009](#); [DeHaan et al., 2015](#); [Ferracuti and Lind, 2025](#)), we construct two complementary measures that incorporate both the volume and economic significance of concurrent earnings announcements.

The first measure, $\log(\text{Size_EA})$, is the natural logarithm of the value-weighted count of earnings announcements on a given day, where each announcement is weighted by the announcing firm’s market capitalization relative to the total market capitalization of all firms announcing earnings during the same year-quarter. This weighting scheme ensures that the measure captures not only the number of concurrent announcements but also their aggregate economic importance, which is an important determinant of investors’ information acquisition intensity and capital-market outcomes ([Frederickson and Zolotoy, 2016](#); [Driskill et al., 2020](#)). The second measure, *Size_EA_intensity*, is the year-quarter decile

rank of the value-weighted daily announcement count, rescaled to the $(0, 1)$ interval so that a one-unit change represents a shift from the bottom to the top decile within the quarter. This transformation preserves cross-day variation in announcement intensity while reducing the influence of outlier days and facilitating economic interpretation of the estimated coefficients.

Both measures are designed to capture the aggregate interpretation burden facing market participants on a given day. Days with numerous, economically significant earnings announcements impose greater demands on investors' capacity to process complex financial disclosures in real time—precisely the conditions under which the integration-cost advantage of generative AI should be most salient.

4.3 Sample Formation and Descriptive Statistics

Our sample period spans August 2024 through October 2025, a window that captures the rapid adoption phase of consumer-facing generative AI platforms while providing sufficient variation in earnings-announcement intensity across multiple quarterly reporting cycles. The empirical analysis draws on three distinct samples, each defined by the data requirements of a particular set of tests.

The primary web-traffic sample consists of 4,548 website-day observations constructed from daily SimilarWeb traffic data for the fifteen platforms in our classification (ten generative AI platforms and five traditional search engines) over the sample period. This sample is used for the baseline tests, cross-sectional analyses, and robustness tests. The referral-traffic sample tracks monthly referral visits from each source platform to downstream financial websites, yielding 1,303 source-destination-month observations. Because SimilarWeb reports referral data only at the monthly frequency, this sample operates at a coarser time resolution than the web-traffic sample and is used exclusively for the mechanism test.

The firm-level sample, used for the market-consequence analysis, contains 14,923

firm-quarter observations obtained by merging data from multiple standard sources. Stock trading data from CRSP are used to construct two event-window market outcomes: *Illiquidity*, defined using the [Amihud \(2002\)](#) illiquidity ratio, and *Bid-Ask Spread*, defined using quoted bid-ask spreads. Both outcomes are measured over the earnings-announcement trading day and the following trading day; for announcements released after market close, the window begins on the next trading day. Firm financial data from Compustat provide the basis for a comprehensive set of controls, including firm size, book-to-market ratio, return on assets, leverage, cash holdings, capital intensity, sales growth, intangible asset intensity, and financing activities. Analyst coverage is obtained from I/B/E/S. We further obtain retail trading data from TAQ, social media attention data from Context Analytics, and earnings conference call transcript length from S&P Capital IQ. All continuous variables are winsorized at the 1st and 99th percentiles.

Table 1 presents summary statistics for the key variables across all three samples. In the web-traffic sample, the mean of log daily visits is 14.69 with a standard deviation of 2.74, and the interquartile range spans from 13.11 to 16.95, reflecting substantial scale dispersion between the largest platforms and smaller generative AI entrants. Generative AI platforms constitute approximately 65 percent of the sample observations, consistent with the larger number of GenAI domains in the treatment group. The earnings-announcement measures display meaningful cross-day variation: $\log(\text{Size_EA})$ has a mean of -6.29 and a standard deviation of 2.98, while *Size_EA_intensity* has a mean of 0.50 and a standard deviation of 0.32. This variation is essential for the identification strategy because our estimates rely on the differential traffic response of generative AI platforms to fluctuations in announcement intensity.

In the referral-traffic sample, the mean scaled referral traffic is 2.42 with a standard deviation of 5.60, and the median is 0.33, indicating that referral intensity is right-skewed and that a relatively small number of source-destination pairs account for disproportionately large referral flows. Approximately 54 percent of destination financial

websites in the sample provide free access, allowing us to distinguish referrals to broadly accessible financial portals from referrals to subscription-based destinations.

In the firm-level sample, the mean Amihud illiquidity ratio is 0.13 with a standard deviation of 3.96, and the mean bid-ask spread is 0.49 with a standard deviation of 1.13. The distributions are right-skewed, consistent with meaningful variation in liquidity across firms around earnings announcements. Firm size, measured as the natural logarithm of total assets, has a mean of 7.29 and a standard deviation of 2.26, and average analyst coverage is 7.39 analysts per firm. These magnitudes indicate that the sample spans a broad cross-section of the earnings-announcement population, which is important for the generalizability of the cross-sectional and market-consequence analyses.

5 Research Design and Empirical Results

5.1 Baseline Evidence: Earnings Announcements and the Reallocation toward Generative AI

We begin by examining whether investor attention shifts from traditional search engines toward generative AI platforms when earnings-related information-processing demand rises. The central prediction is that, if GenAI lowers the cost of synthesizing and interpreting dense financial information, its relative usage should increase on days with more concentrated earnings announcements. We test this prediction by estimating the following specification:

$$\log(\text{Web Traffic})_{i,t} = \beta_0 + \beta_1 \text{GenAI}_i \times \text{EA}_t + \gamma' \mathbf{X}_{i,t} + \mu_i + \eta_t + \varepsilon_{i,t} \quad (1)$$

where $\log(\text{Web Traffic})_{i,t}$ is the natural logarithm of daily visits to website i on day t , GenAI_i is an indicator equal to one for generative AI platforms and zero for traditional search engines, and EA_t is the earnings-announcement intensity on day t . The coefficient of

interest, β_1 , captures the differential traffic response of generative AI platforms relative to traditional search engines during periods of high earnings-announcement intensity. The vector $\mathbf{X}_{i,t}$ contains the six visitor-characteristic controls described above. Website fixed effects μ_i absorb time-invariant differences across platforms, while time fixed effects η_t control for common daily shocks to internet usage. We estimate four specifications that vary the EA intensity measure ($\log(\text{Size_EA})$ or Size_EA_intensity) and the granularity of time fixed effects (year-quarter or year-date). Standard errors are clustered at the domain level throughout.

Table 2 reports results. Across all four specifications, the coefficient on $\text{GenAI} \times \text{EA}$ is positive and statistically significant. In Column (1), which uses $\log(\text{Size_EA})$ as the intensity measure and includes domain and year-quarter fixed effects, the estimated coefficient is 0.016 ($t = 4.143$); Column (2) replaces the continuous measure with the decile-rank index Size_EA_intensity and yields 0.140 ($t = 3.047$). Columns (3) and (4) tighten the design by substituting year-date fixed effects for year-quarter fixed effects, thereby absorbing all common daily variation in web traffic, and produce similar estimates of 0.017 ($t = 4.343$) and 0.147 ($t = 3.259$). Economically, these coefficients imply that a one-unit increase in log value-weighted earnings-announcement intensity is associated with approximately a 1.6 to 1.7 percent increase in relative traffic to generative AI platforms; moving from the lowest to the highest decile of announcement intensity corresponds to a roughly 14.0 to 14.7 percent larger increase in generative AI traffic relative to traditional search. These magnitudes are economically meaningful given the scale of the platforms involved and the daily frequency of website visits. The adjusted R^2 is large in every specification, confirming that the fixed-effect structure accounts for the vast majority of variation in daily platform traffic and that the remaining identifying variation comes from the interaction between platform type and earnings-announcement intensity.

Another notable feature of the results is the coefficient on EA itself—the main effect of earnings-announcement intensity, which captures the traffic response of traditional

search engines (the base group when GenAI equals zero). In Columns (1) and (2), this coefficient is -0.001 ($t = -0.407$) and -0.010 ($t = -0.496$), respectively, and is statistically indistinguishable from zero. This null finding stands in sharp contrast to a large body of prior work showing that traditional search engine activity rises significantly during earnings-announcement periods. [Da et al. \(2011\)](#) establish Google Search Volume Index as a direct measure of retail investor attention that increases around value-relevant events, and [Drake et al. \(2012\)](#) document that abnormal Google search volume spikes at the earnings announcement date. The absence of a positive EA effect for traditional search in our sample period—a period in which generative AI has become widely accessible—suggests that the historical role of search engines in earnings-related information acquisition has been attenuated by the availability of generative AI as an alternative information-processing technology.

Taken together, the significantly positive $GenAI \times EA$ interaction and the null EA main effect paint a coherent picture of user migration: investor information demand during earnings announcements has not diminished, but has shifted toward synthesis-oriented platforms. This pattern is difficult to reconcile with a view in which generative AI traffic simply rises with general online activity or with a secular increase in AI adoption unrelated to financial information processing, because the time fixed effects absorb common daily shocks and the website fixed effects control for platform-specific scale. Rather, the evidence suggests that investors disproportionately migrate toward generative AI platforms when they must rapidly interpret a large volume of earnings news—precisely the setting where the integration-cost advantage of synthesis-oriented platforms is most salient.

5.2 Cross-Sectional Heterogeneity

The baseline results establish that generative AI platforms attract disproportionately more traffic when earnings-announcement intensity is high. A natural follow-up question is whether this reallocation is uniform across all types of earnings events or concentrated

in settings where GenAI's synthesis function is likely to be more valuable. We examine this question along two related dimensions. First, we test whether the GenAI response is stronger when earnings information is more difficult to interpret. Second, motivated by the policy concern about broad access to financial information, we test whether the response is stronger when the announcing firms attract greater retail attention. We therefore examine three moderators: the number of business segments, the magnitude of the earnings surprise, and retail attention.

We estimate a triple-interaction specification that allows the GenAI response to high earnings-announcement intensity to vary with these announcement-day characteristics. For each earnings-announcement day, we compute the average characteristic among announcing firms and set the moderator indicator to one when this daily average is in the top quartile of its sample distribution.

Business Segments captures organizational complexity: firms with more business segments are more difficult to evaluate because investors must integrate information across more lines of business. *Earnings Surprise* captures the extent to which the announcement conveys unexpected information. *Retail Attention* captures the degree to which the announcing firms are traded by retail investors. The coefficient of interest is the triple interaction $GenAI \times EA \times Moderator$, which identifies whether the differential GenAI response on high-announcement days is stronger when the information environment is more demanding to process.

Table 3 reports the results consistent with this prediction. The coefficient on the triple interaction is positive and statistically significant for all three moderators. The coefficient is 0.121 for Business Segments, 0.121 for Earnings Surprise, and 0.127 for Retail Attention, with t -statistics of 2.248, 1.879, and 2.320, respectively.¹ The first two results indicate that the migration toward GenAI is stronger when earnings information is more demanding to interpret: investors must synthesize information across more complex business operations

¹The sample is slightly smaller for Earnings Surprise ($N = 4,237$) because the measure requires I/B/E/S analyst consensus forecasts, which are unavailable for some announcement days.

or evaluate larger deviations from expectations. The retail-attention result provides a complementary insight. It indicates that GenAI-based information acquisition is more pronounced when announcing firms attract greater retail trading interest, consistent with GenAI being particularly salient in settings where broad investor participation is more important.

Taken together, these patterns sharpen the interpretation of the baseline findings. They are difficult to reconcile with alternative explanations that attribute the GenAI traffic surge to generic curiosity or non-financial usage. The reallocation toward GenAI is strongest when the earnings news is harder to interpret and when the announcing firms are more retail-relevant. This pattern is consistent with GenAI reducing the cost of accessing synthesized interpretations in earnings-announcement settings. Whether this increased use improves the information environment, however, depends on how effectively investors process and act on AI-assisted information, a question we examine through the market-outcome tests later.

5.3 Mechanism: Referral Traffic to Financial Websites

The preceding analyses show that generative AI traffic rises on earnings-announcement days and that the effect concentrates where interpretation demands and retail relevance are highest. A remaining question is whether this traffic increase genuinely reflects financial-information processing or merely captures unrelated platform usage that coincides with the earnings calendar. We address this concern by tracing the downstream destinations to which generative AI platforms direct their users. If investors use GenAI to process earnings information, a portion of their sessions should culminate in visits to primary financial data sources, producing an observable referral pathway from GenAI platforms to financial websites. Moreover, because generative AI platforms generally cannot access or synthesize content behind paywalls, we expect this referral effect to concentrate among freely accessible financial portals. We test these predictions by

estimating the following specification:

$$Referral\ Traffic_{i,j,t} = \beta_0 + \beta_1 GenAI_i \times EA_t \times Free_j + \gamma' \mathbf{X}_{i,t} + \mu_{i,j} + \eta_t + \varepsilon_{i,j,t} \quad (2)$$

where $Referral\ Traffic_{i,j,t}$ is the monthly referral traffic from source platform i to financial website j in month t , scaled by the monthly average web traffic of the platform and multiplied by 1,000. The destination websites include major financial information portals such as Yahoo Finance, Seeking Alpha, and Bloomberg. $Free_j$ is an indicator equal to one if the financial website provides free access without a subscription and zero otherwise. The triple interaction $GenAI \times EA \times Free$ identifies whether the differential referral response of generative AI platforms during high-announcement months is concentrated among freely accessible financial portals—the destinations most likely to serve as downstream complements for investors processing earnings disclosures with an AI-generated synthesis. Because SimilarWeb reports referral data only at the monthly level, this analysis is conducted at the source-destination-month frequency ($N = 1,303$). We include source-destination pair fixed effects and year-month fixed effects; lower-order interaction terms are included but omitted from the presentation for brevity.

Table 4 presents the results. The coefficient on $GenAI \times EA \times Free$ is positive and statistically significant across both EA intensity measures: 0.196 ($t = 2.860$) using $\log(Size\ EA)$ and 0.102 ($t = 2.296$) using $Size\ EA\ intensity$. These estimates indicate that during months with more intensive earnings-announcement activity, generative AI platforms generate significantly more referral traffic to free financial-information websites relative to both subscription-based financial destinations and the referral traffic generated by traditional search engines.

This finding is informative on two dimensions. First, it links the observed GenAI traffic surge to financially relevant downstream activity rather than to broad, non-financial browsing, directly mitigating the identification concern that motivated this test. Second, the concentration among free-access destinations is economically intuitive: investors

who use GenAI to synthesize earnings information are likely to seek verification or supplementary detail from easily accessible primary sources, and the inability of many GenAI models to access paywalled content makes free portals the natural complement. Because free financial websites are more accessible to nonprofessional investors than paywalled institutional platforms, this pattern is also consistent with GenAI-supported information acquisition occurring through channels available to a broad investor base. Together, the referral evidence supports the interpretation that generative AI platforms serve as an entry point for financial-information acquisition during earnings periods—a gateway through which investors synthesize complex disclosures before turning to primary data sources for confirmation.

5.4 Market Consequences

The evidence thus far establishes that investors reallocate attention from traditional search to generative AI during earnings announcements, that this reallocation intensifies with interpretation demands and retail participation, and that it generates measurable downstream traffic to financial information websites. We next examine whether this shift in information acquisition is associated with changes in the capital-market information environment. The prediction is theoretically ambiguous. If generative AI lowers the cost of interpreting public disclosures for a broader set of investors, liquidity should improve because more investors can process earnings information in a timely manner and adverse selection should decline. By contrast, if generative AI primarily enhances the ability of already-informed or institutionally sophisticated investors to extract and trade on earnings news, adverse selection could increase, leading to lower liquidity and wider spreads.

We conduct the analysis at the firm level. Because web traffic is observed at the platform-day level, we relate daily generative AI and traditional search traffic to market outcomes for firms announcing earnings on each day. We focus on three firm characteristics that capture settings in which AI-assisted information processing is likely

to matter for market quality. *Retail Attention* identifies firms with greater retail trading interest, where reductions in processing costs may be especially relevant for broad investor participation. *Tweet Volume* captures the intensity of public discussion around the firm, where investors face a greater need to organize and interpret dispersed information signals. *Conference Call Length* captures the volume of the firm’s earnings-related disclosure, where the synthesis function of generative AI should be particularly useful. We estimate the following specification:

$$Y_{i,t} = \beta_0 + \beta_1 \text{GenAI Traffic}_t \times EA_t \times \text{Characteristic}_{i,t} \\ + \beta_2 \text{TradSearch Traffic}_t \times EA_t \times \text{Characteristic}_{i,t} + \gamma' \mathbf{X}_{i,t} + \mu_i + \eta_t + \varepsilon_{i,t} \quad (3)$$

where $Y_{i,t}$ is either illiquidity (Amihud ratio) or bid-ask spread for firm i around its earnings announcement. Both outcomes are computed over the earnings-announcement trading day and the following trading day, with timing adjusted for after-market announcements. GenAI Traffic_t and $\text{TradSearch Traffic}_t$ are the natural logarithms of daily visits to generative AI platforms and traditional search engines, respectively. EA_t is an indicator for days with above-median earnings-announcement intensity. $\text{Characteristic}_{i,t}$ equals one when the corresponding firm characteristic is in the top sample quartile. The inclusion of both generative AI and traditional search traffic interactions within the same specification allows us to test whether any market effect is specific to generative AI or instead reflects a broader increase in online information-seeking. Firm fixed effects, year-quarter fixed effects, and a comprehensive set of firm-level controls are included in all specifications. Lower-order interaction terms are included but omitted from the tabulated results for brevity.

Table 5 reports the results. For illiquidity, the coefficient on $\text{GenAI Traffic} \times EA \times \text{Characteristic}$ is negative and statistically significant across all three firm characteristics: -0.054 ($t = -4.551$) for Retail Attention, -0.063 ($t = -3.851$) for Tweet Volume, and

-0.057 ($t = -4.726$) for Conference Call Length. For bid-ask spread, the corresponding coefficients are also negative and significant: -0.104 ($t = -3.372$), -0.128 ($t = -2.903$), and -0.088 ($t = -2.499$), respectively.² In contrast, the corresponding traditional search interactions are positive and significant or marginally significant across all specifications. For illiquidity, the coefficients on $TradSearch\ Traffic \times EA \times Characteristic$ are 0.045, 0.057, and 0.049; for bid-ask spread, the coefficients are 0.069, 0.100, and 0.080.

The joint pattern of lower illiquidity and narrower bid-ask spreads associated with generative AI traffic, but not with traditional search traffic, suggests an improvement in the earnings-announcement information environment. Lower illiquidity and narrower spreads indicate lower trading frictions and reduced adverse selection, consistent with faster incorporation of disclosure content into prices. The effects are concentrated among firms with greater retail trading interest, more public discussion, and longer earnings conference calls—settings in which investors are likely to benefit more from tools that organize and synthesize complex information. This pattern is difficult to attribute to generic online information-seeking. Instead, it is consistent with GenAI improving market quality by reducing the cost of interpreting earnings disclosures when those disclosures are most salient or demanding to process.

6 Robustness Tests

We subject the baseline findings to a battery of robustness tests that address alternative explanations and measurement concerns. Each test targets a specific threat to the interpretation of the main results, and collectively they strengthen the case that the documented reallocation toward generative AI reflects event-specific financial-information processing rather than confounding patterns in platform usage.

²The sample is smaller when Tweet Volume is the partitioning variable ($N = 11,335$) because Context Analytics data coverage does not span our full sample period.

6.1 Placebo Tests Using Non-Earnings Corporate Events

An important concern is that the baseline results may reflect a broader tendency for GenAI usage to rise around major corporate news events, rather than a mechanism tied specifically to the interpretation of earnings disclosures. We address this concern by replacing earnings-announcement intensity with the intensity of three major non-earnings corporate events: mergers and acquisitions, legal actions, and executive changes. Table 6 reports two sets of tests. Columns (1)–(3) include concurrent earnings-announcement intensity and its interaction with GenAI; Columns (4)–(6) exclude these earnings controls.

The results indicate that the positive GenAI response is specific to earnings announcements. In Columns (1)–(3), the interactions between GenAI and non-earnings event intensity are statistically insignificant for M&A events, legal actions, and executive changes: 0.045 ($t = 1.195$), -0.012 ($t = -0.403$), and 0.006 ($t = 0.132$), respectively. By contrast, the interaction between GenAI and earnings-announcement intensity remains positive and significant in each specification. Columns (4)–(6) show that, when earnings intensity is omitted, the non-earnings event interactions become positive and significant. This pattern suggests that non-earnings corporate events partly coincide with periods of elevated earnings-related information demand. Once concurrent earnings-announcement intensity is included, these events do not generate a distinct GenAI response. The evidence therefore supports the interpretation that the documented reallocation toward GenAI is tied to the processing of earnings information rather than to corporate news more generally.

6.2 Introduction of Google AI Mode

The preceding results suggest that investors migrate toward GenAI because these platforms offer synthesis capabilities that traditional search engines historically lacked. A direct implication is that, if traditional search becomes more synthesis-oriented, the relative GenAI response during earnings periods should diminish. We test this implication

by exploiting Google’s introduction of AI Mode, which embedded generative synthesis capabilities directly into Google Search and made the dominant traditional search engine more similar to standalone GenAI platforms.

Table 7 augments the baseline specification with the interaction $GenAI \times EA \times PostAIMode$, where *PostAIMode* equals one for observations on or after March 1, 2025, and *EA* equals one for days with above-median earnings-announcement intensity. The coefficient on this triple interaction is negative in both specifications: -0.096 ($t = -1.685$) using $\log(Size_EA)$ and -0.113 ($t = -1.713$) using *Size_EA_intensity*. The negative sign indicates that, after Google introduces AI Mode, the relative increase in standalone GenAI traffic on high-announcement days becomes smaller. At the same time, the coefficient on $GenAI \times EA$ remains positive and significant, indicating that standalone GenAI platforms continue to attract greater traffic during high-announcement periods. Together, these results are consistent with a mechanism rooted in synthesis functionality. When search engines incorporate generative synthesis, the functional distinction between the two platform types narrows and the incremental traffic response of standalone GenAI platforms attenuates.

6.3 Alternative Measures and Timing Choices

We next examine whether the baseline results are sensitive to alternative measures of platform usage and earnings-announcement intensity. The baseline analysis uses total website visits as the dependent variable. One concern is that the results may be driven by more intensive usage among a fixed group of users rather than by a broader shift in user activity. Table 8 addresses this concern by replacing total visits with unique visitors. The interaction between GenAI and earnings-announcement intensity remains positive and statistically significant: 0.015 ($t = 3.788$) using $\log(Size_EA)$ and 0.127 ($t = 2.667$) using *Size_EA_intensity*. This evidence indicates that the earnings-related increase in GenAI usage is not limited to repeated visits by an existing user base; it also reflects a broader

increase in unique users.

Table 9 examines whether the results depend on the value-weighted construction of earnings-announcement intensity. We replace the baseline measures with the raw number of firms announcing earnings on a given day, using both $\log(\text{Num_EA})$ and Num_EA_intensity . The interaction between GenAI and earnings-announcement intensity remains positive and significant in both specifications: 0.031 ($t = 2.509$) and 0.135 ($t = 3.176$), respectively. Thus, the documented migration toward GenAI does not depend on weighting announcements by firm size.

We also consider the timing of information availability. Earnings announcements released after the market close may not be fully incorporated into trading until the next trading day. Table 10 therefore constructs timing-aware measures of earnings-announcement intensity. Columns (1)–(2) measure EA intensity using announcements released before or during market hours on the announcement day. Columns (3)–(4) assign after-market announcements to the next trading day, when the information first becomes available for trading. The $\text{GenAI} \times \text{EA}$ interaction remains positive and significant across both timing conventions. The coefficients are 0.019 ($t = 3.258$) and 0.153 ($t = 2.411$) in the pre/during-market specification, and 0.013 ($t = 2.221$) and 0.121 ($t = 1.996$) in the information-available specification. These results indicate that the baseline findings are not an artifact of how announcement dates are assigned; the GenAI response continues to align with the timing of earnings-information availability.

6.4 Timely Search Functionality

During our sample period, some GenAI platforms introduced timely web-search functionality. This development raises a distinct concern. The paper’s interpretation emphasizes GenAI’s ability to reduce integration costs by synthesizing financial information. However, if the earnings-announcement response is concentrated among

GenAI platforms with timely search capability, the observed traffic increase could partly reflect improved retrieval of current information rather than synthesis alone. We examine this possibility by introducing *TimelySearch*, an indicator equal to one when a platform provides timely search functionality. Traditional search engines are coded as having timely search throughout the sample.

Table 11 reports the results. The coefficient on $GenAI \times EA \times TimelySearch$ is positive and statistically significant or marginally significant: 0.009 ($t = 2.022$) using $\log(Size_EA)$ and 0.082 ($t = 1.775$) using *Size_EA_intensity*. These estimates indicate that timely search functionality makes GenAI platforms more attractive during earnings-announcement periods. Importantly, the lower-order interaction $GenAI \times EA$ remains positive and significant in both specifications, with coefficients of 0.009 ($t = 5.041$) and 0.074 ($t = 4.013$). Thus, while timely search features strengthen the GenAI response, they do not fully explain the baseline result. The evidence suggests that the earnings-announcement effect is not solely driven by GenAI platforms becoming better retrieval tools; rather, the traffic response remains present after accounting for the availability of timely search functionality.

6.5 User Composition Around Earnings Announcements

Finally, we examine whether the composition of platform users changes around periods of high earnings-announcement intensity. In the main specifications, monthly user characteristics are included as controls to absorb variation in platform user bases. Table 12 instead examines these characteristics. This analysis asks whether the relative increase in GenAI traffic during earnings periods is accompanied by a shift toward users with greater investment-related interests.

The results provide corroborating evidence that the GenAI response reflects financial-information processing. The coefficient on $GenAI \times EA$ is positive and significant when the dependent variable is Investing Overlap: 0.070 ($t = 1.981$). In contrast, the corresponding coefficients are small and insignificant for News Media Overlap, Science

Overlap, and Social Network Overlap. This pattern indicates that, during periods of more intensive earnings announcements, GenAI platforms attract relatively more users who also visit investment-related websites, but not users associated with other broad interest categories. Together with the referral-traffic evidence in Table 4, these results strengthen the interpretation that the increase in GenAI usage around earnings announcements reflects demand for financial information rather than a general shift in online activity.

7 Conclusion

This paper studies how generative AI affects investor information acquisition around earnings announcements. Using daily web-traffic data for generative AI platforms and traditional search engines, we find that generative AI usage increases disproportionately on high-earnings-announcement days, while traditional search traffic does not respond to announcement intensity. The reallocation concentrates where interpretation demands and retail relevance are greatest—among earnings events involving more complex firms, larger earnings surprises, and higher retail attention—and produces measurable downstream referral traffic to financial information websites. At the firm level, generative AI traffic during high-announcement periods is associated with lower illiquidity and narrower bid-ask spreads for firms with high retail attention, high social media attention, and longer conference call transcripts; traditional search traffic shows no analogous association. Placebo tests confirm that the effect is specific to earnings announcements rather than general corporate news, and the introduction of Google’s AI Mode attenuates the differential demand for standalone generative AI platforms, consistent with a mechanism rooted in synthesis functionality.

These findings have implications for both research and policy. Search activity has long served as the primary digital trace of investor attention, especially for retail investors, but our evidence suggests that this measure may now understate earnings-related information

acquisition when investors increasingly turn to synthesis-oriented tools. More broadly, the findings speak to a central concern in disclosure regulation: public information reduces asymmetry only if investors can use it. GenAI may function as a user-side interpretation technology that makes complex earnings disclosures more accessible to a broader investor base. At the same time, its benefits may be uneven if more sophisticated investors have greater awareness, better access, stronger prompting ability, and richer complementary information. The evidence that GenAI activity is strongest in retail-attention settings and is associated with lower illiquidity and narrower spreads is consistent with an improved information environment. It highlights why AI-assisted interpretation is relevant to ongoing debates over investor protection, technology-mediated investing, and the democratization of financial markets.

Our analysis has limitations. SimilarWeb data capture observable visits to websites and applications but do not include API-based usage or interactions embedded within third-party tools, which may constitute a growing share of AI-assisted information processing. Our identification strategy relies on variation in earnings-announcement intensity across days and cannot directly observe the content of individual queries or the identity of individual users. As the boundary between retrieval and interpretation continues to blur—through the integration of generative capabilities into search engines and the expansion of AI tools into professional workflows—understanding which investors use GenAI, how effectively they use it, and how this technological shift affects price discovery and market efficiency remains an important direction for future research.

Table 1. Summary Statistics

Variable	N	Mean	SD	p25	p50	p75
<i>Website Traffic</i>						
Web Traffic	4,548	14.69	2.74	13.11	14.19	16.95
GenAI	4,548	0.65	0.48	0.00	1.00	1.00
PostAIMode	4,548	0.57	0.49	0.00	1.00	1.00
TimelySearch	4,548	0.85	0.36	1.00	1.00	1.00
<i>Earnings Announcement Measures</i>						
log(Size_EA)	4,548	-6.29	2.98	-7.99	-5.62	-3.98
Size_EA Intensity	4,548	0.50	0.32	0.22	0.44	0.78
<i>Referral Traffic</i>						
Referral Traffic	1,303	2.42	5.60	0.01	0.33	1.83
Free	1,303	0.54	0.50	0.00	1.00	1.00
<i>Market Outcome</i>						
Illiquidity	14,923	0.13	3.96	0.00	0.00	0.02
Bid-Ask Spread	14,923	0.49	1.13	0.05	0.13	0.41
<i>Website User Characteristics</i>						
Female Share	4,548	40.48	4.86	36.96	41.80	43.71
Young User Share	4,548	69.34	5.83	67.43	70.64	73.60
Investing Overlap	4,548	0.83	0.42	0.57	0.76	1.08
News Media Overlap	4,548	2.25	1.23	1.43	1.99	2.55
Science Overlap	4,548	1.02	0.51	0.54	0.99	1.35
Social Network Overlap	4,548	16.17	5.42	12.27	15.63	18.31
<i>Firm-Quarter Controls</i>						
Firm Size	14,923	7.29	2.26	5.90	7.41	8.82
Book-to-Market	14,923	0.48	1.67	0.16	0.39	0.78
ROA	14,923	-0.02	0.16	-0.02	0.00	0.02
Leverage	14,923	0.34	0.30	0.12	0.30	0.47
Cash	14,923	0.14	0.16	0.03	0.08	0.18
Capital Intensity	14,923	0.23	0.24	0.05	0.13	0.33
Sales Growth	14,923	0.30	6.17	-0.04	0.02	0.09
Intangibles	14,923	0.02	0.03	0.00	0.00	0.02
Financing	14,923	0.05	0.14	0.00	0.00	0.04
Analyst Coverage	14,923	7.39	6.69	2.00	5.00	11.00

Notes: This table presents the descriptive statistics for the main variables used in the analyses over the sample period from August 2024 to October 2025. All variables are defined in Appendix A. All continuous variables are winsorized at the 1% and 99% levels.

Table 2. Earnings Announcement Intensity and User Migration to Generative AI

Dependent Variable EA Measure	(1)	(2)	(3)	(4)
	Web Traffic			
	$\log(\text{Size_EA})$	Size_EA_Intensity	$\log(\text{Size_EA})$	Size_EA_Intensity
GenAI $\times$ EA	0.016*** (4.143)	0.140*** (3.047)	0.017*** (4.343)	0.147*** (3.259)
EA	-0.001 (-0.407)	-0.010 (-0.496)		
Female Share	0.313*** (6.117)	0.313*** (6.129)	0.316*** (6.869)	0.316*** (6.880)
Young User Share	0.204*** (3.119)	0.204*** (3.126)	0.198*** (3.118)	0.198*** (3.123)
Investing Overlap	-0.716*** (-2.854)	-0.716*** (-2.832)	-0.863*** (-3.476)	-0.864*** (-3.457)
News Media Overlap	0.176 (1.216)	0.176 (1.217)	0.159 (0.985)	0.159 (0.984)
Science Overlap	-0.565 (-1.229)	-0.565 (-1.223)	-0.449 (-0.955)	-0.450 (-0.954)
Social Network Overlap	-0.041 (-0.737)	-0.042 (-0.744)	-0.052 (-0.911)	-0.052 (-0.915)
N	4,548	4,548	4,548	4,548
Adj. R ²	0.977	0.977	0.977	0.977
Domain FE	Yes	Yes	Yes	Yes
Year-Quarter FE	Yes	Yes	No	No
Year-Date FE	No	No	Yes	Yes

Notes: This table reports how earnings announcement (EA) intensity affects web traffic to generative AI platforms relative to traditional search engines. The dependent variable *Web Traffic* is the natural logarithm of daily website visits. *GenAI* equals one for Generative AI platforms and zero for traditional search engines. EA intensity is measured using two proxies: $\log(\text{Size_EA})$, the natural logarithm of the value-weighted count of daily earnings announcements, and Size_EA_Intensity , the year-quarter decile rank of the value-weighted count (rescaled to the (0, 1) interval). Columns (1)–(2) include domain and year-quarter fixed effects; columns (3)–(4) replace year-quarter with date fixed effects. All variables are defined in Appendix A. The t-statistics are reported in parentheses and are clustered at the domain level. ***, **, and * denote statistical significance at the 1%, 5%, and 10% levels, respectively.

Table 3. Heterogeneity in User Migration to Generative AI in Response to Earnings

	(1)	(2)	(3)
Dependent Variable	Web Traffic		
Moderator	Business Segments	Earnings Surprise	Retail Attention
GenAI $\times$ EA $\times$ Moderator	0.121** (2.248)	0.121* (1.879)	0.127** (2.320)
GenAI $\times$ EA	0.037 (1.107)	0.019 (0.724)	0.050 (1.196)
GenAI $\times$ Moderator	-0.078* (-2.036)	-0.012 (-0.285)	-0.101 (-1.309)
EA $\times$ Moderator	-0.070* (-2.006)	-0.044 (-1.106)	-0.084** (-2.534)
EA	0.023 (1.610)	0.023 (1.750)	0.019 (1.006)
Moderator	0.033 (1.386)	-0.014 (-0.941)	0.062* (1.939)
N	4,548	4,237	4,548
Adj. R ²	0.977	0.977	0.977
Controls	Yes	Yes	Yes
Domain FE	Yes	Yes	Yes
Year-Quarter FE	Yes	Yes	Yes

Notes: This table reports cross-sectional heterogeneity in the effect of earnings announcement intensity on web traffic to generative AI platforms relative to traditional search engines. *EA* equals one when *EA* intensity is above the sample median. For each earnings announcement day, we compute the average of the relevant characteristic among announcing firms and define *Moderator* as one when the characteristic exceeds the sample top quartile. Columns (1)–(3) use *Business Segments*, *Earnings Surprise*, and *Retail Attention* as moderators, respectively. *EA* intensity is measured using $\log(\text{Size} \cdot \text{EA})$. All variables are defined in Appendix A. The t-statistics are reported in parentheses and are clustered at the domain level. ***, **, and * denote statistical significance at the 1%, 5%, and 10% levels, respectively.

Table 4. Referral Traffic Patterns: Generative AI Platforms vs. Traditional Search

	(1)	(2)
Dependent Variable	Referral Traffic	
EA Measure	$\log(\text{Size_EA})$	Size_EA_Intensity
GenAI $\times$ EA $\times$ Free	0.196*** (2.860)	0.102** (2.296)
Female Share	0.118 (0.275)	0.119 (0.278)
Young User Share	0.112 (0.421)	0.112 (0.423)
Investing Overlap	−1.891 (−0.912)	−1.887 (−0.909)
News Media Overlap	1.672 (1.608)	1.670 (1.605)
Science Overlap	−1.375 (−0.504)	−1.380 (−0.504)
Social Network Overlap	0.660* (1.961)	0.661* (1.957)
N	1,303	1,303
Adj. R ²	0.798	0.798
Lower-Order Interactions Included	Yes	Yes
Controls	Yes	Yes
Source-Destination Pair FE	Yes	Yes
Year-Month FE	Yes	Yes

Notes: This table examines referral traffic from generative AI platforms and traditional search engines to financial websites, and how it varies with monthly EA intensity and destination website access policy. The dependent variable is *Referral Traffic*, defined as monthly referral visits from a source domain (either a generative AI platform or traditional search engine) to financial websites, scaled by the average daily web traffic of the source domain within the month. *Free* equals one if the destination financial website offers free content access without a subscription. *EA* denotes aggregate monthly earnings announcement intensity, measured using $\log(\text{Size_EA})$ and Size_EA_Intensity . All variables are defined in Appendix A. The t-statistics are reported in parentheses and are clustered at the source domain level. ***, **, and * denote statistical significance at the 1%, 5%, and 10% levels, respectively.

Table 5. Market Outcomes — Firm Level: Illiquidity and Bid-Ask Spread

Dependent Variables	(1)	(2)	(3)	(4)	(5)	(6)
	Illiquidity			Bid-Ask Spread		
Character	Retail Attention	Tweet Volume	Conference Call Length	Retail Attention	Tweet Volume	Conference Call Length
GenAI Traffic×EA×Character	−0.054*** (−4.551)	−0.063*** (−3.851)	−0.057*** (−4.726)	−0.104*** (−3.372)	−0.128*** (−2.903)	−0.088** (−2.499)
TradSearch Traffic×EA×Character	0.045*** (2.976)	0.057** (2.070)	0.049*** (3.402)	0.069* (1.858)	0.100* (1.855)	0.080** (2.115)
GenAI Traffic×EA	0.063*** (5.557)	0.057*** (3.667)	0.063*** (5.657)	0.114*** (3.763)	0.100*** (2.590)	0.110*** (3.743)
TradSearch Traffic×EA	−0.055*** (−3.670)	−0.052* (−1.869)	−0.054*** (−3.932)	−0.081** (−2.218)	−0.065 (−1.259)	−0.085** (−2.507)
N	14,923	11,335	14,923	14,923	11,335	14,923
Adj. R ²	0.710	0.727	0.710	0.812	0.828	0.812
Lower-Order Interactions Included	Yes	Yes	Yes	Yes	Yes	Yes
Controls	Yes	Yes	Yes	Yes	Yes	Yes
Firm FE	Yes	Yes	Yes	Yes	Yes	Yes
Year-Quarter FE	Yes	Yes	Yes	Yes	Yes	Yes

Notes: This table examines whether the user-migration response to earnings announcements is associated with firm-level market outcomes. Firm-level market outcomes are measured by *Illiquidity* (Amihud ratio) and *Bid-Ask Spread*. Because web traffic is observed only at the platform level, we allocate daily platform traffic across firms announcing earnings on that day using the following firm characteristics: *Retail Attention*, *Tweet Volume*, and *Conference Call Length*. *Character* equals one if the corresponding firm characteristic exceeds the top sample quartile, and zero otherwise. *EA* equals one for days with above-median EA intensity, and zero otherwise. EA intensity is measured using $\log(\text{Size} \cdot \text{EA})$. *GenAI Traffic* and *TradSearch Traffic* are the natural logarithms of daily visits to generative AI platforms and traditional search engines, respectively. Control variables include firm size, book-to-market, ROA, leverage, cash holdings, capital intensity, sales growth, intangible intensity, financing activities, and analyst coverage. All variables are defined in Appendix A. The t-statistics are reported in parentheses and are clustered at the firm level. ***, **, and * denote statistical significance at the 1%, 5%, and 10% levels, respectively.

Table 6. Placebo Tests: Non-Earnings Corporate Events

	(1)	(2)	(3)	(4)	(5)	(6)
Dependent Variables	Web Traffic					
Event Type	M&A	Legal	Executive	M&A	Legal	Executive
GenAI $\times$ Event Intensity	0.045 (1.195)	-0.012 (-0.403)	0.006 (0.132)	0.113*** (2.871)	0.065** (2.072)	0.101** (2.347)
Event Intensity	0.082*** (3.397)	0.075*** (3.536)	0.091*** (3.251)	0.092*** (2.850)	0.086*** (3.058)	0.097*** (2.635)
EA Intensity	0.021 (0.869)	0.022 (0.802)	0.009 (0.427)			
GenAI $\times$ EA Intensity	0.148*** (3.816)	0.171*** (4.227)	0.163*** (3.693)			
Female Share	0.327*** (6.289)	0.327*** (6.297)	0.328*** (6.290)	0.328*** (6.264)	0.328*** (6.276)	0.328*** (6.261)
Young User Share	0.205*** (3.053)	0.204*** (3.048)	0.204*** (3.049)	0.206*** (3.042)	0.204*** (3.031)	0.205*** (3.033)
Investing Overlap	-0.583** (-2.412)	-0.585** (-2.434)	-0.582** (-2.405)	-0.572** (-2.368)	-0.576** (-2.392)	-0.570** (-2.359)
News Media Overlap	0.146 (0.961)	0.145 (0.962)	0.146 (0.963)	0.147 (0.963)	0.146 (0.962)	0.147 (0.967)
Science Overlap	-0.285 (-0.551)	-0.287 (-0.554)	-0.285 (-0.551)	-0.280 (-0.548)	-0.284 (-0.553)	-0.280 (-0.547)
Social Network Overlap	-0.050 (-0.873)	-0.051 (-0.889)	-0.051 (-0.884)	-0.050 (-0.858)	-0.051 (-0.888)	-0.051 (-0.877)
N	6,526	6,526	6,526	6,526	6,526	6,526
Adj. R ²	0.975	0.975	0.975	0.975	0.975	0.975
Domain FE	Yes	Yes	Yes	Yes	Yes	Yes
Year-Quarter FE	Yes	Yes	Yes	Yes	Yes	Yes

Notes: This table presents placebo tests examining whether the intensity of non-earnings corporate events differentially affects web traffic to generative AI platforms versus traditional search engines. Placebo event intensity is the year-quarter decile rank of the value-weighted count of M&A announcements, legal-actions disclosures, or executive-change announcements on a given day. Columns (1)–(3) control for *EA Intensity* measured using *Size.EA Intensity*. Columns (4)–(6) exclude these EA controls. All variables are defined in Appendix A. The t-statistics are reported in parentheses and are clustered at the domain level. ***, **, and * denote statistical significance at the 1%, 5%, and 10% levels, respectively.

Table 7. The Introduction of Google AI Mode and the AI–Search Traffic Differential

	(1)	(2)
Dependent Variable	Web Traffic	
EA Measure	$\log(\text{Size_EA})$	Size_EA_Intensity
GenAI $\times$ EA $\times$ PostAIMode	−0.096* (−1.685)	−0.113* (−1.713)
GenAI $\times$ EA	0.146*** (2.764)	0.170*** (2.596)
GenAI $\times$ PostAIMode	0.433* (1.824)	0.436* (1.836)
EA $\times$ PostAIMode	0.051* (1.853)	0.063* (1.666)
EA	−0.048** (−2.468)	−0.064** (−2.230)
PostAIMode	−0.257* (−1.765)	−0.259* (−1.776)
Female Share	0.298*** (5.480)	0.298*** (5.500)
Young User Share	0.216*** (3.325)	0.216*** (3.332)
Investing Overlap	−0.687*** (−2.627)	−0.688*** (−2.626)
News Media Overlap	0.215 (1.391)	0.215 (1.390)
Science Overlap	−0.470 (−1.086)	−0.471 (−1.085)
Social Network Overlap	−0.034 (−0.593)	−0.034 (−0.592)
N	4,548	4,548
Adj. R ²	0.978	0.978
Domain FE	Yes	Yes
Year-Quarter FE	Yes	Yes

Notes: This table examines whether the introduction of Google AI Mode changed the traffic differential between Generative AI platforms and traditional search engines on high earnings-announcement days. *PostAIMode* equals one for periods on or after March 1, 2025 and zero otherwise. *EA* equals one for days with above-median (within year-quarter) EA intensity, and zero otherwise. EA intensity is measured using $\log(\text{Size_EA})$ and Size_EA_Intensity . Control variables include the average demographic characteristics of website users. All variables are defined in Appendix A. The t-statistics are reported in parentheses and are clustered at the domain level. ***, **, and * denote statistical significance at the 1%, 5%, and 10% levels, respectively.

Table 8. Alternative Traffic Measure: Unique Visitors

	(1)	(2)
Dependent Variable	Unique Visitors	
EA Measure	$\log(\text{Size_EA})$	Size_EA_Intensity
GenAI $\times$ EA	0.015*** (3.788)	0.127*** (2.667)
EA	−0.001 (−0.648)	−0.011 (−0.546)
Female Share	0.293*** (6.039)	0.294*** (6.049)
Young User Share	0.171*** (2.982)	0.171*** (2.989)
Investing Overlap	−0.580** (−2.312)	−0.580** (−2.294)
News Media Overlap	0.120 (0.918)	0.120 (0.919)
Science Overlap	−0.624 (−1.380)	−0.624 (−1.375)
Social Network Overlap	−0.040 (−0.795)	−0.040 (−0.803)
N	4,548	4,548
Adj. R ²	0.975	0.975
Domain FE	Yes	Yes
Year-Quarter FE	Yes	Yes

Notes: This table examines how earnings announcement (EA) intensity affects unique visitors to generative AI platforms relative to traditional search engines. The dependent variable is the natural logarithm of daily website unique visitors. *EA* denotes EA intensity, measured using two proxies: $\log(\text{Size_EA})$, the natural logarithm of the value-weighted count of daily earnings announcements, and Size_EA_Intensity , the year-quarter decile rank of the value-weighted count, rescaled to the (0, 1) interval. All variables are defined in Appendix A. t-statistics are reported in parentheses and clustered at the domain level. ***, **, and * denote statistical significance at the 1%, 5%, and 10% levels, respectively.

Table 9. Alternative Earnings Announcement Intensity Measures

	(1)	(2)
Dependent Variable	Web Traffic	
EA Measure	$\log(\text{Num_EA})$	Num_EA_Intensity
GenAI $\times$ EA	0.031** (2.509)	0.135*** (3.176)
EA	−0.001 (−0.116)	0.026 (0.199)
Female Share	0.313*** (6.097)	0.313*** (6.098)
Young User Share	0.206*** (3.125)	0.205*** (3.122)
Investing Overlap	−0.709*** (−2.824)	−0.712*** (−2.821)
News Media Overlap	0.175 (1.209)	0.175 (1.205)
Science Overlap	−0.573 (−1.232)	−0.568 (−1.221)
Social Network Overlap	−0.040 (−0.705)	−0.040 (−0.702)
N	4,548	4,548
Adj. R ²	0.977	0.977
Domain FE	Yes	Yes
Year-Quarter FE	Yes	Yes

Notes: This table replicates the baseline results (Table 2) using alternative measures of earnings announcement (EA) intensity. $\log(\text{Num_EA})$ is the natural logarithm of the number of firms announcing earnings on a given day. Num_EA_Intensity is the corresponding year-quarter decile rank (rescaled to the (0, 1) interval). All variables are defined in Appendix A. The t-statistics are reported in parentheses and are clustered at the domain level. ***, **, and * denote statistical significance at the 1%, 5%, and 10% levels, respectively.

Table 10. Timing-Aware Earnings Announcement Intensity and Web Traffic

	(1)	(2)	(3)	(4)
Dependent Variable	Web Traffic			
Timing family	Pre/During		Info Available	
EA Measure	log(Size_EA)	Size_EA_Intensity	log(Size_EA)	Size_EA_Intensity
GenAI $\times$ EA	0.019*** (3.258)	0.153** (2.411)	0.013** (2.221)	0.121** (1.996)
EA	-0.005 (-1.449)	-0.038 (-1.179)	-0.007*** (-2.596)	-0.055** (-2.283)
Female Share	0.314*** (6.064)	0.314*** (6.071)	0.312*** (6.039)	0.312*** (6.048)
Young User Share	0.207*** (3.104)	0.207*** (3.111)	0.207*** (3.102)	0.207*** (3.109)
Investing Overlap	-0.735*** (-2.882)	-0.737*** (-2.859)	-0.726*** (-2.812)	-0.729*** (-2.803)
News Media Overlap	0.172 (1.162)	0.171 (1.159)	0.174 (1.173)	0.173 (1.169)
Science Overlap	-0.570 (-1.237)	-0.573 (-1.239)	-0.566 (-1.234)	-0.566 (-1.232)
Social Network Overlap	-0.038 (-0.662)	-0.038 (-0.671)	-0.036 (-0.636)	-0.037 (-0.641)
N	4,331	4,331	4,418	4,418
Adj. R ²	0.977	0.977	0.977	0.977
Domain FE	Yes	Yes	Yes	Yes
Year-Quarter FE	Yes	Yes	Yes	Yes

Notes: This table examines whether the relation between earnings announcement intensity and web traffic depends on the timing of information availability. Columns (1)–(2) measure EA intensity using announcements released before or during market hours on the announcement day (*Pre/During*). Columns (3)–(4) use the same construction but assign after-market announcements to the next trading day, when the information first becomes available for trading (*Info Available*). EA intensity is measured as $\log(\text{Size_EA})$. All variables are defined in Appendix A. The t-statistics are reported in parentheses and are clustered at the domain level. ***, **, and * denote statistical significance at the 1%, 5%, and 10% levels, respectively.

Table 11. Timely Search Feature and the Earnings Announcement Effect

Dependent Variable EA Measure	(1)	(2)
	Web Traffic log(Size_EA)	Size_EA_Intensity
GenAI $\times$ EA $\times$ TimelySearch	0.009** (2.022)	0.082* (1.775)
GenAI $\times$ EA	0.009*** (5.041)	0.074*** (4.013)
GenAI $\times$ TimelySearch	0.923 (0.723)	0.828 (0.670)
EA $\times$ TimelySearch	0.000 (0.201)	0.003 (0.131)
EA	-0.000 (-0.130)	-0.006 (-0.237)
TimelySearch	-0.409 (-0.349)	-0.421 (-0.357)
Female Share	0.297*** (6.844)	0.298*** (6.849)
Young User Share	0.192*** (3.367)	0.192*** (3.375)
Investing Overlap	-0.627*** (-2.880)	-0.628*** (-2.849)
News Media Overlap	0.140 (0.924)	0.140 (0.924)
Science Overlap	-0.500 (-1.136)	-0.499 (-1.128)
Social Network Overlap	-0.040 (-0.800)	-0.040 (-0.808)
N	4,548	4,548
Adj. R ²	0.978	0.978
Controls	Yes	Yes
Domain FE	Yes	Yes
Year-Quarter FE	Yes	Yes

Notes: This table examines whether timely search features affect the earnings-announcement traffic response. *TimelySearch* equals one when a platform provides timely search functionality, and zero otherwise. It equals one for traditional search engines throughout the sample. EA intensity is measured using $\log(\text{Size_EA})$ and *Size_EA_Intensity*. All variables are defined in Appendix A. The t-statistics are reported in parentheses and are clustered at the domain level. ***, **, and * denote statistical significance at the 1%, 5%, and 10% levels, respectively.

Table 12. Changes in User Characteristics Around Earnings Announcement Intensity

	(1)	(2)	(3)	(4)
Dependent Variable	Investing	News Media	Science	Social Network
GenAI $\times$ EA	0.070** (1.981)	-0.032 (-0.699)	0.004 (0.296)	-0.186 (-1.153)
EA	-0.003 (-0.140)	0.037** (2.251)	0.021*** (3.540)	0.073** (2.268)
N	215	215	215	215
Adj. R ²	0.607	0.898	0.883	0.815
Domain FE	Yes	Yes	Yes	Yes
Year-Quarter FE	Yes	Yes	Yes	Yes

Notes: This table examines whether earnings announcements shift the user composition of Generative AI platforms relative to traditional search engines. The dependent variables are the monthly percentage of users overlapping with investing, news media, science, and social network domains. EA intensity is measured using the natural logarithm of monthly average value-weighted count of daily earnings announcements. All variables are defined in Appendix A. The t-statistics are reported in parentheses and are clustered at the domain level. ***, **, and * denote statistical significance at the 1%, 5%, and 10% levels, respectively.

Appendix A Variable Definitions

Variable	Definition	Source
<i>Website Traffic</i>		
Web Traffic	Natural logarithm of daily U.S. website visits to a generative AI platform or traditional search platform.	SimilarWeb
GenAI	Indicator equal to 1 for generative AI platforms (ChatGPT, Perplexity, Claude, Gemini, Copilot, Grok, Deepseek, Meta AI, Poe, and Qwen) and 0 for traditional search engines (Google, Bing, Yahoo, DuckDuckGo, and Quora).	Authors' classification
PostAIMode	Indicator equal to one for observations on or after March 1, 2025, when Google introduced its AI Mode feature, and 0 otherwise.	Google
TimelySearch	Indicator equal to one when a platform provides a timely search feature. It equals one for traditional search engines throughout the sample.	Authors' classification
<i>Earnings Announcement Intensity Measures</i>		
log(Size_EA)	Natural logarithm of the value-weighted count of daily earnings announcements, where each announcement is weighted by the announcing firm's market capitalization relative to the total market capitalization of all firms announcing earnings during the same year-quarter.	I/B/E/S, Compustat, and Authors' calculation
Size_EA_Intensity	Year-quarter decile rank of the value-weighted count of daily earnings announcements, where each announcement is weighted by the announcing firm's market capitalization relative to the total market capitalization of all firms announcing earnings during the same year-quarter. The rank is rescaled to the (0,1) interval.	I/B/E/S, Compustat, and Authors' calculation
<i>Referral Variables</i>		
Referral Traffic	Monthly referrals from a source platform to a financial destination website, scaled by the average daily web traffic of the source platform within the month ($\times 1000$).	SimilarWeb
Free	Indicator equal to one if the destination website provides free access without a subscription requirement.	Authors' classification
<i>Market Outcomes and Firm Controls</i>		
Illiquidity	Average Amihud illiquidity over the earnings-announcement trading day and the following trading day, where daily Amihud illiquidity is defined as the absolute value of the daily stock return divided by dollar trading volume. For announcements released after market close, the window begins on the next trading day t ($\times 10^6$).	CRSP

Continued on next page

(Appendix A continued)

Variable	Definition	Source
Bid-Ask Spread	Average quoted bid-ask spread over the earnings-announcement trading day and the following trading day, where the daily spread is defined as the ask price minus the bid price, divided by the midpoint of the bid and ask prices. For announcements released after market close, the window begins on the next trading day t ($\times 100$).	CRSP
Website User Characteristics		
Female Share	Monthly percentage of website users who identify as female.	SimilarWeb
Young User Share	Monthly percentage of website users aged 18–34.	SimilarWeb
Investing Overlap	Monthly percentage of website users who also visit investment-related websites.	SimilarWeb
News Media Overlap	Monthly percentage of website users who also visit news and media websites.	SimilarWeb
Science Overlap	Monthly percentage of website users who also visit science or other educational websites.	SimilarWeb
Social Network Overlap	Monthly percentage of website users who also visit social-network websites.	SimilarWeb
Partition Variables		
Business Segments	Number of business segments of the firm in the fiscal quarter whose earnings are announced.	Compustat
Earnings Surprise	Difference between actual quarterly earnings per share and the mean analyst forecast for the firm in the fiscal quarter whose earnings are announced.	I/B/E/S
Retail Attention	Average fraction of all retail trading volume that goes to the stock in the previous calendar quarter.	TAQ
Tweet Volume	Average tweet volume related to the firm in the previous calendar quarter.	Context Analytics
Conference Call Length	Number of words in the firm's earnings conference call transcripts.	S&P Capital IQ
Firm-Quarter Controls		
Firm Size	Natural logarithm of total book assets for the firm in the fiscal quarter whose earnings are announced.	Compustat
Book-to-Market	Book-to-market equity ratio for the firm in the fiscal quarter whose earnings are announced.	Compustat
ROA	Return on assets for the firm in the fiscal quarter whose earnings are announced, calculated as net income divided by total assets.	Compustat
Leverage	Leverage ratio for the firm in the fiscal quarter whose earnings are announced, calculated as total debt divided by total assets.	Compustat
Cash	Cash and short-term investments scaled by total assets for the firm in the fiscal quarter whose earnings are announced.	Compustat

Continued on next page

(Appendix A continued)

Variable	Definition	Source
Capital Intensity	Capital expenditures scaled by total assets for the firm in the fiscal quarter whose earnings are announced.	Compustat
Sales Growth	Year-over-year change in net sales scaled by lagged assets for the firm in the fiscal quarter whose earnings are announced.	Compustat
Intangibles	Intangible assets scaled by total assets for the firm in the fiscal quarter whose earnings are announced.	Compustat
Financing	Equity or long-term debt issuance scaled by assets for the firm in the fiscal quarter whose earnings are announced.	Compustat
Analyst Coverage	Number of analysts following the firm in the fiscal quarter whose earnings are announced.	I/B/E/S

References

- Amihud, Y. (2002). Illiquidity and stock returns: Cross-section and time-series effects. *Journal of Financial Markets*, 5(1):31–56.
- Barber, B. M., Huang, X., Odean, T., and Schwarz, C. (2022). Attention-induced trading and returns: Evidence from Robinhood users. *Journal of Finance*, 77(6):3141–3190.
- Barber, B. M. and Odean, T. (2000). Trading is hazardous to your wealth: The common stock investment performance of individual investors. *Journal of Finance*, 55(2):773–806.
- Barber, B. M. and Odean, T. (2008). All that glitters: The effect of attention and news on the buying behavior of individual and institutional investors. *Review of Financial Studies*, 21(2):785–818.
- Bertomeu, J., Lin, Y., Liu, Y., and Ni, Z. (2025). The impact of generative AI on information processing: Evidence from the ban of ChatGPT in Italy. *Journal of Accounting and Economics*, 80:101782.
- Bhattacharya, N., Cho, Y. J., and Kim, J. B. (2018). Leveling the playing field between large and small institutions: Evidence from the SEC’s XBRL mandate. *The Accounting Review*, 93(5):51–71.
- Blankespoor, E., Croom, J., and Grant, S. M. (2026a). Generative AI and investor processing of financial information. *Journal of Accounting Research*, 64.
- Blankespoor, E., deHaan, E., and Li, Q. (2026b). Generative AI in financial reporting. *Journal of Accounting Research*, 64.
- Blankespoor, E., deHaan, E., and Marinovic, I. (2020). Disclosure processing costs, investors’ information choice, and equity market outcomes: A review. *Journal of Accounting and Economics*, 70(2–3):101344.
- Blankespoor, E., deHaan, E., Wertz, J., and Zhu, C. (2019). Why do individual investors disregard accounting information? The roles of information awareness and acquisition costs. *Journal of Accounting Research*, 57(1):53–84.
- Bloomberg–SimilarWeb (2026). SimilarWeb and Bloomberg expand data available on the Bloomberg Terminal. Business Wire, February 2026.
- Bradshaw, M. T., Ma, C., Yost, B., and Zou, Y. (2025). Generative AI use by capital market information intermediaries: Evidence from Seeking Alpha. Working paper.
- Cheng, Q., Lin, P., and Zhao, Y. (2025). Does generative AI facilitate investor trading? Early

- evidence from ChatGPT outages. *Journal of Accounting and Economics*, 80:101821.
- Croom, J. (2025). Interactivity and illusions of ability: How using generative AI affects investor judgments. *Journal of Accounting Research*, 63.
- Da, Z., Engelberg, J., and Gao, P. (2011). In search of attention. *Journal of Finance*, 66(5):1461–1499.
- DeHaan, E., Lawrence, A., and Litjens, R. (2025). AI-generated content in the financial press. *Management Science*.
- DeHaan, E., Shevlin, T., and Thornock, J. (2015). Market (in)attention and the strategic scheduling and timing of earnings announcements. *Journal of Accounting and Economics*, 60(1):36–55.
- Drake, M. S., Roulstone, D. T., and Thornock, J. R. (2012). Investor information demand: Evidence from Google searches around earnings announcements. *Journal of Accounting Research*, 50(4):1001–1040.
- Driskill, M., Kirk, M. P., and Tucker, J. W. (2020). Concurrent earnings announcements and analysts' information production. *The Accounting Review*, 95(2):165–189.
- Ecker, F., Li, X., Li, Y., and Wu, F. (2025). How stock market participants use generative artificial intelligence: Evidence from user-platform interaction data. Working paper.
- Ferracuti, E. and Lind, G. (2025). Macroeconomic information acquisition around earnings clusters. *Journal of Accounting Research*, 63.
- Frederickson, J. R. and Zolotoy, L. (2016). Competing earnings announcements: Which announcement do investors process first? *The Accounting Review*, 91(2):441–462.
- Hirshleifer, D., Lim, S. S., and Teoh, S. H. (2009). Driven to distraction: Extraneous events and underreaction to earnings news. *The Journal of Finance*, 64(5):2289–2325.
- Hodge, F. D., Kennedy, J. J., and Maines, L. A. (2004). Does search-facilitating technology improve the transparency of financial reporting? *The Accounting Review*, 79(3):687–703.
- Jackson, R. J. (2019). Statement on reducing investor protections around private markets. SEC Commissioner Statement, December 18.
- Lawrence, A., Ryans, J., Sun, E., and Laptev, N. (2018). Earnings announcement promotions: A Yahoo Finance field experiment. *Journal of Accounting and Economics*, 66(2–3):399–414.
- Liu, Y., Huang, A., and Wang, Y. (2025). The global rise of generative AI: Trends, adoption,

- and market dynamics. Working paper.
- Melumad, S. and Yun, H.-K. (2025). Experimental evidence of the effects of large language models versus web search on depth of learning. *PNAS Nexus*, 4:pgaf316.
- Pew Research Center (2025). How Americans view and use AI-powered search. Technical report, Pew Research Center, Washington, DC.
- SEC (1998). A plain English handbook: How to create clear SEC disclosure documents. Office of Investor Education and Assistance.
- SEC (2021). Request for information and comments on broker-dealer and investment adviser digital engagement practices. Release No. 34-92766.
- SEC (2023). Conflicts of interest associated with the use of predictive data analytics by broker-dealers and investment advisers. Release Nos. 34-97990; IA-6353.
- Semrush (2025a). Investigating ChatGPT search: Insights from 80 million clickstream records. Semrush Blog, January.
- Semrush (2025b). SearchGPT vs. Google and Bing: A quick comparison of search engines. Semrush Blog, January.
- SimilarWeb (2024a). Data-as-a-service: Websites dataset. Available at: <https://www.similarweb.com/corp/daas/datasets/websites/>.
- SimilarWeb (2024b). Our data. Available at: <https://www.similarweb.com/corp/ourdata/>.
- SimilarWeb (2024c). SimilarWeb data methodology. Available at: <https://support.similarweb.com/hc/en-us/articles/360001631538>.
- Spatharioti, S. E., Rothschild, D. M., Goldstein, D. G., and Hofman, J. M. (2025). Comparing traditional and LLM-based search for consumer choice. In *Proceedings of CHI 2025*. ACM.
- Zubair, R., Alhassan, I., and Bello, G. (2025). User experience in generative AI versus traditional search engines. *Publications*, 13:14.